

Assignment 2

Ghatta Trivedi

2025-02-07

R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
summary(cars)
```

```
##      speed      dist
##  Min.   : 4.0    Min.   :  2.00
##  1st Qu.:12.0    1st Qu.: 26.00
##  Median :15.0    Median : 36.00
##  Mean   :15.4    Mean   : 42.98
##  3rd Qu.:19.0    3rd Qu.: 56.00
##  Max.   :25.0    Max.   :120.00
```

Including Plots

You can also embed plots, for example:



Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.

```
source("C:\\Users\\trive\\Documents\\Applied Stat Learning R\\codeday6.pck")
codeday6.pck
```

```
## [1] "codeday6.pck"      "my.bootstrap1"     "gui.bootstrap1"    "my.bootstrap2"
## [5] "gui.bootstrap2"    "my.bootstrap3"     "gui.bootstrapxy"   "my.dat.plot5"
## [9] "my.dat.plot5a"
```

```
NOAA <- read.csv("C:\\Users\\trive\\Documents\\Applied Stat Learning R\\NOAA+GISS 2024.csv")
NOAA
```

```
##   year X.disaster delta.temp
## 1 1980          3      0.27
## 2 1981          1      0.33
## 3 1982          3      0.13
## 4 1983          5      0.30
## 5 1984          2      0.16
## 6 1985          5      0.12
## 7 1986          2      0.19
## 8 1987          0      0.33
## 9 1988          1      0.41
##10 1989          5      0.29
##11 1990          3      0.44
```

## 12 1991	4	0.43
## 13 1992	6	0.23
## 14 1993	4	0.24
## 15 1994	6	0.32
## 16 1995	4	0.45
## 17 1996	4	0.35
## 18 1997	3	0.48
## 19 1998	9	0.63
## 20 1999	5	0.42
## 21 2000	2	0.42
## 22 2001	2	0.54
## 23 2002	4	0.63
## 24 2003	7	0.62
## 25 2004	5	0.55
## 26 2005	5	0.69
## 27 2006	6	0.64
## 28 2007	5	0.66
## 29 2008	11	0.54
## 30 2009	7	0.65
## 31 2010	5	0.73
## 32 2011	16	0.61
## 33 2012	11	0.64
## 34 2013	9	0.66
## 35 2014	8	0.75
## 36 2015	10	0.90
## 37 2016	15	1.02
## 38 2017	16	0.93
## 39 2018	14	0.85
## 40 2019	14	0.99
## 41 2020	22	1.02
## 42 2021	20	0.85
## 43 2022	18	0.89
## 44 2023	28	1.17
## 45 2024	26	1.28

#commented out code

to fullfill the first section of the assignment, we have to make a function that takes in a data frame
`first <- function(mat=NOAA,i=3,j=2,zxlab,zylab,zmain,zcol,do.sqrt=F,nboot=10000){`
#function(mat=NOAA,i=3,j=2,zxlab,zylab,zmain,zcol,do.sqrt=F,do.plot=T,in.boot=T)
`#stat.out<-list(smstrmod=smstr,smstrlin=smstr.lin,resid.mod=resid1,resid.lin=resid2,dfmod=dfmod,dflin=d`

#first we need to set up the plot area to show plots side by side
`par(mfrow=c(1,2))`

#this line calls for the function my.dat.plot5, which will: plot the data with a smoothing spline and
the argument do.plot = T means it will show the plot and in.boot = F means it's not in the bootstrap
`stat.out0<-my.dat.plot5(mat,i,j,zxlab,zylab,zmain,zcol,do.sqrt,do.plot=T,in.boot=F)`

#now, we need to store results
this one stores the fstat
`F0<-stat.out0$F`

#gets the residual from the linear fit
`resid0<-stat.out0$resid.lin`

```

#initialize a vecotre to store f stats from the bootstrapped fits
bootvec<-NULL

#get predicted values from the linear model using the first column of the data
y0<-predict(stat.out0$smstrlin,mat[,i])$y

#next we need to create a copy of the input matrix for bootstrapping
matb<-mat

#now we need to loop to perform bootstrap iterations
for(i1 in 1:nboot){

  #lets us know progress by printing 500,1000,etc
  if(floor(i1/500)==(i1/500)){print(i1)}

  #sample residuals with replacement to create a bootstrapped residual vector
  residb<-sample(resid0,replace=T)

  #add the bootstrapped residuals to the predicted values to generate new bootstrapped rensponse va
  Yb<-y0+residb

  # now we can replace the dependent variable with the bootstrapped response vals
  matb[,j]<-Yb

  #again call the my.dat.plot5 again but this time with in.boot = T so it doesn't plot anything onl
  stat.outb<-my.dat.plot5(matb,i,j,zxlab,zylab,zmain,zcol,do.sqrt,do.plot=F,in.boot=T)
  #store the f stat from the bootstrapped model fit
  bootvec<-c(bootvec,stat.outb$F)
}

# now, outside the loop we will calculate the pval by determining the proportion of bootstrapped f-st
pvalboot<-sum(bootvec>F0)/nboot
#create the boxplot of the bootstrapped f-stat
boxplot(bootvec)
#add the bootstrap p val to the lis tof output from the original fit
stat.out0$pvalboot<-pvalboot
#returns the full results, including hte original fit stat and the bootstrap model's p-val
stat.out0
}



#basically this function is performing a model comparison between a smoothing spline and a linear fit.

first(NOAA, "X.disaster", "delta.temp", "Year", "Delta Temp", "Disaster vs Temp", 1, do.sqrt = FALSE, nl

## [1] 500
## [1] 1000

```

Disaster vs Temp



```
## $smstrmod
## Call:
## smooth.spline(x = mat[, i], y = mat[, j])
##
## Smoothing Parameter spar= 1.493422 lambda= 7188.814 (26 iterations)
## Equivalent Degrees of Freedom (Df): 2.000011
## Penalized Criterion (RSS): 0.3294866
## GCV: 0.02673542
##
## $smstrlin
## Call:
## smooth.spline(x = mat[, i], y = mat[, j], df = 2)
##
## Smoothing Parameter spar= 1.49992 lambda= 8007.977 (32 iterations)
## Equivalent Degrees of Freedom (Df): 2.00001
## Penalized Criterion (RSS): 0.3294866
## GCV: 0.02673542
##
## $resid.mod
## [1] -0.1253104918  0.0051411772 -0.2653104918 -0.1657621628 -0.2000846568
## [6] -0.3457621628 -0.1700846568  0.0403670109  0.0851411772 -0.1757621628
## [11]  0.0446895082 -0.0005363276 -0.2709879954 -0.1905363276 -0.1809879954
## [16]  0.0194636724 -0.0805363276  0.0846895082  0.0233345467 -0.0457621628
## [21]  0.0599153432  0.1799153432  0.1994636724  0.0837861766  0.0842378372
## [26]  0.2242378372  0.1390120046  0.1942378372 -0.1371170364  0.1137861766
## [31]  0.2642378372 -0.2432457759 -0.0371170364  0.0533345467  0.1785603564
```

```
## [36] 0.2581087489 0.2019799475 0.0767542241 0.0672056823 0.2072056823
## [41] -0.0445999761 -0.1441485930 -0.0336971964 -0.1059540797 0.0744972827
##
## $resid.lin
## [1] -0.1253105089 0.0051411456 -0.2653105089 -0.1657621652 -0.2000846812
## [6] -0.3457621652 -0.1700846812 0.0403669721 0.0851411456 -0.1757621652
## [11] 0.0446894911 -0.0005363373 -0.2709879908 -0.1905363373 -0.1809879908
## [16] 0.0194636627 -0.0805363373 0.0846894911 0.0233345685 -0.0457621652
## [21] 0.0599153188 0.1799153188 0.1994636627 0.0837861878 0.0842378348
## [26] 0.2242378348 0.1390120092 0.1942378348 -0.1371170089 0.1137861878
## [31] 0.2642378348 -0.2432457564 -0.0371170089 0.0533345685 0.1785603735
## [36] 0.2581087741 0.2019799710 0.0767542436 0.0672057088 0.2072057088
## [41] -0.0445999956 -0.1441485979 -0.0336971879 -0.1059541481 0.0744972311
##
## $dfmod
## [1] 2.000011
##
## $dflin
## [1] 2
##
## $F
## [1] 0.1733031
##
## $P
## [1] 7.32035e-05
##
## $n
## [1] 45
##
## $pvalboot
## [1] 0.42
```

```
my.bootstrap2 <-function(mat=NOAA,i=3,j=2,zxlab,zylab,zmain,zcol,do.sqrt=F,nboot=10000,pred.bound=T,con
#function(mat=NOAA,i=3,j=2,zxlab,zylab,zmain,zcol,do.sqrt=F,do.plot=T,in.boot=T)
#stat.out<-list(smstrmod=smstr,smstrlin=smstr.lin,resid.mod=resid1,resid.lin=resid2,dfmod=dfmod,dflin

# again, we need to set up the plot so:
par(mfrow=c(1,1))

#Calling the my.dat.plot5 to plot the data with a smoothing spline adn linear fit and to get the resi
stat.out0<-my.dat.plot5(mat,i,j,zxlab,zylab,zmain,zcol,do.sqrt,do.plot=T,in.boot=F)

#now, we Save residuals
resid0<-stat.out0$resid.mod

#Create an empty and initialize it with 0s matrix to save results of bootstrapped data
bootmat<-NULL

# get the predicted values from the non-linear model
y0<-predict(stat.out0$smstrmod,mat[,i])$y

#create a copy of the original dataset for bootstrapping
matb<-mat
```

```

# in this for loop we will perform the bootstrapping by resampling residuals with replacement
#so we have make this loop fo 10000 times
for(i1 in 1:nboot){
  #print progress every 500 iterations
  if(floor(i1/500)==(i1/500)){print(i1)}

  #sample residuals with replacemnt
  residb<-sample(resid0,replace=T)

  #generate bootstrapped results by adding the bootstrapped residuals
  Yb<-y0+residb

  #Replace the dependent variable in matb with the bootstrapped values
  matb[,j]<-Yb
  #print(matb)

  #call the my.dat.plot5 to fit the model on the bootstrapped data without plotting
  stat.outb<-my.dat.plot5(matb,i,j,xzlab,zylab,zmain,zcol,do.sqrt,do.plot=F,in.boot=T)

  #extract residuals and predicted values from the bootstrapped model fit
  Ybp<-predict(stat.outb$smstrmod,matb[,i])$y

  # construct confidencce intervals based on the choice of 'pivotal' and 'pred.bound'
  if(pred.bound){
    if(pivotal){
      bootmat<-rbind(bootmat,stat.outb$resid.mod+Ybp-y0)
    }else{
      bootmat<-rbind(bootmat,stat.outb$resid.mod+Ybp)
    }
  }else{
    if(pivotal){
      bootmat<-rbind(bootmat,Ybp-y0)
    }else{
      bootmat<-rbind(bootmat,Ybp)
    }
  }
}

#calculate confidence intervals for the bootstrapped data
alpha<-(1-conf.lev)/2
my.quant<-function(x,a=alpha){quantile(x,c(a,1-a))}
bounds<-apply(bootmat,2,my.quant)

#adjust bounds for pivotal bootstrap
if(pivotal){
  bounds[1,]<-y0-bounds[1,]
  bounds[2,]<-y0-bounds[2,]
}

#sort the data based on the x calues
x<-mat[,i]
if(do.sqrt){
  x<-sqrt(x)
}

```

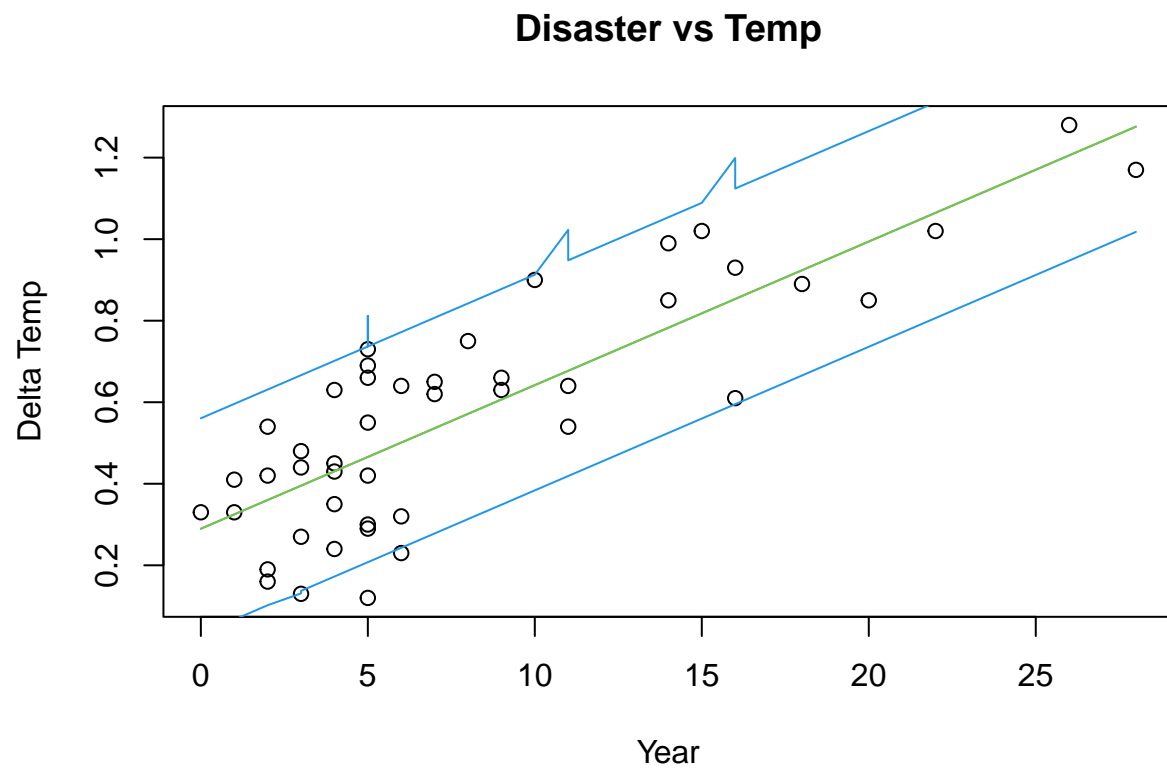
```

}
o1<-order(x)

#add confidence interval lines to the plot
lines(x[o1],bounds[1,o1],col=zcol+2)
lines(x[o1],bounds[2,o1],col=zcol+2)
}

my.bootstrap2(NOAA, "X.disaster", "delta.temp", "Year", "Delta Temp", "Disaster vs Temp", 2)

```



```

## [1] 500
## [1] 1000
## [1] 1500
## [1] 2000
## [1] 2500
## [1] 3000
## [1] 3500
## [1] 4000
## [1] 4500
## [1] 5000
## [1] 5500
## [1] 6000
## [1] 6500
## [1] 7000
## [1] 7500

```



```
## [1] 8000
## [1] 8500
## [1] 9000
## [1] 9500
## [1] 10000
```

```
my.bootstrap3 <-
function(mat=NOAA,i=3,j=2,zxlab,zylab,zmain,zcol,do.sqrt=F,nboot=10000,pred.bound=T,conf.lev=.95,pivotal=
  #set up the plot again
  par(mfrow=c(1,1))

  #call the plot function to plot data
  stat.out0<-my.dat.plot5(mat,i,j,zxlab,zylab,zmain,zcol,do.sqrt,do.plot=T,in.boot=F)

  #intialize matrix to store bootstrapped data again
  bootmat<-NULL

  #get the predicted values from the plot again
  y0<-predict(stat.out0$smstrmod,mat[,i])$y

  #create a copy of the original dataset for bootstrapping
  matb<-mat
  #get the length of the dataset
  nm<-length(matb[,1])

  #perform xy bootstrapping by sampling entire rows of the data
  for(i1 in 1:nboot){
    #checking progress again every 500 times
    if(floor(i1/500)==(i1/500)){print(i1)}

    # sample rows with replacement ie entire rows are resampled, not just residuals
    zed<-sample(nm,replace=T)
    matb<-mat[zed,]

    #print(matb)
    #fit the model with the bootstrapped data
    stat.outb<-my.dat.plot5a(matb,mat,i,j,zxlab,zylab,zmain,zcol,do.sqrt,do.plot=F,in.boot=T)

    #get predicted values from the bootstrapped model fit
    Ybp<-predict(stat.outb$smstrmod,mat[,i])$y

    #constuct confidence intervals based on the choice
    if(pred.bound){
      if(pivotal){
        #pivotal bootstrap adjusts for the baseline prediction qwhich is to subtract y0 to create c
        bootmat<-rbind(bootmat,stat.outb$resid.mod+Ybp-y0)
      }else{
        #non pivotal directly adds the residuals to the predictions
        bootmat<-rbind(bootmat,stat.outb$resid.mod+Ybp)
      }
    }else{
      if(pivotal){
        #pivotal bootstrap with no boundary would ajust residuals + predictions from bootstrap
```

```

        bootmat<-rbind(bootmat,Ybp-y0)
      }else{
        #no boundary for non pivotal bootstrap
        bootmat<-rbind(bootmat,Ybp)
      }
    }
  }

  #calculate quantiles for the bootstrapped data
  #set the alpha level based on ci
  alpha<-(1-conf.lev)/2
  my.quant<-function(x,a=alpha){quantile(x,c(a,1-a))}
  #apply quant function to each colum of the bootstrapped data
  bounds<-apply(bootmat,2,my.quant)

  #if pivotal is used, adjust bounds to account for the baseline model prediction
  if(pivotal){
    #lower bound is adjusted by subtracting y0
    bounds[1,]<-y0-bounds[1,]
    #upper bound is also adjusted by subtracting y0
    bounds[2,]<-y0-bounds[2,]
  }
  # extract x values for plotting
  x<-mat[,i]

  if(do.sqrt){
    #apply square root transformation if required
    x<-sqrt(x)
  }


  #order the data based on the x values
  o1<-order(x)

  #add ci bounds to the plot
  lines(x[o1],bounds[1,o1],col=zcol+2)
  lines(x[o1],bounds[2,o1],col=zcol+2)
}

my.bootstrap3(NOAA, "X.disaster", "delta.temp", "Year", "Delta Temp", "Disaster vs Temp",2)

```

Disaster vs Temp



```
## [1] 500
## [1] 1000
## [1] 1500
## [1] 2000
## [1] 2500
## [1] 3000
## [1] 3500
## [1] 4000
## [1] 4500
## [1] 5000
## [1] 5500
## [1] 6000
## [1] 6500
## [1] 7000
## [1] 7500
## [1] 8000
## [1] 8500
## [1] 9000
## [1] 9500
## [1] 10000
```

```
# commenting on my.datplot5
```

```
my.dat.plot5 <-function(mat=NOAA,i=3,j=2,zxlab,zylab,zmain,zcol,do.sqrt=F,do.plot=T,in.boot=T){
  # if in boot is true then we need to set do sqrt to false. we do this to make sure that sqrt transform
```

```

if(in.boot){
do.sqrt<-F
}

#if dosqrt is dalse no sqroot transformation is applied to the dependent variable
if(!do.sqrt){

  #fit a smoothing spline model to the original dat
  smstr<-smooth.spline(mat[,i],mat[,j])
  #fit a linear smoothing spline because of df=2
  smstr.lin<-smooth.spline(mat[,i],mat[,j],df=2)

  #if do.plot is true, plot the data
  if(do.plot){
    #plot the raw data
    plot(mat[,i],mat[,j],xlab=zxlab,ylab=zylab,main=zmain)
    #add line of non-linear smoothing spline
    lines(smstr,col=zcol)
    #plot the linear smoothing spline
    lines(smstr.lin,col=(zcol+1))
  }

  # compute residuals for the transformed data first one is for the non-linear and the second is for
  resid1<-mat[,j]-predict(smstr,mat[,i])$y
  resid2<-mat[,j]-predict(smstr.lin,mat[,i])$y

}else{
  # fit a smoothing spline model to the original dat
  smstr<-smooth.spline(mat[,i],sqrt(mat[,j]))
  #fit a linear smoothing spline because of df=2
  smstr.lin<-smooth.spline(mat[,i],sqrt(mat[,j]),df=2)
  #if do.plot is true, plot the data
  if(do.plot){
    #plot the raw data
    plot(mat[,i],sqrt(mat[,j]),xlab=zxlab,ylab=zylab,main=zmain)
    #add line of non-linear smoothing spline
    lines(smstr,col=zcol)
    #plot the linear smoothing spline
    lines(smstr.lin,col=(zcol+1))
  }

  # compute residuals for the transformed data first one is for the non-linear and the second is for t
  resid1<-sqrt(mat[,j])-predict(smstr,mat[,i])$y
  resid2<-sqrt(mat[,j])-predict(smstr.lin,mat[,i])$y
}

#get df fpr tje non linear smoothing spline model
dfmod<-smstr$df
#df for the lin model is fixed at 2
dflin<-2

#calculate sum of squares for the eacj of the model
ssmod<-sum(resid1^2) # which is sum of squared residuls
sslin<-sum(resid2^2)

```

```

#compute the difference in sum of squares between the two models
numss<-sslin-ssmod

#get num of data points
n1<-length(mat[,j])

#perform an f test to compare the models
Fstat<-(numss/(dfmod-dflin))/(ssmod/(n1-dfmod))

#calculate the pvalue associated with the fstat
pvalue<-1-pf(Fstat,dfmod-dflin,n1-dfmod)

#return a list contain non linear spline model, linear spline model, residuals for each model, dfs for
stat.out<-list(smstrmod=smstr,smstrlin=smstr.lin,resid.mod=resid1,resid.lin=resid2,dfmod=dfmod,dflin=dflin)
}

```